

# Experiment 7

November 9, 2023

Name: Adwait Purao UID: 2021300101 Batch: B2

```
[ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
```

```
[ ]: df = pd.read_csv("/content/drive/MyDrive/Colab Notebooks/Employee.csv")
```

```
[ ]: df.head()
```

```
[ ]: 
```

	Education	JoiningYear	City	PaymentTier	Age	Gender	EverBenched	\
0	Bachelors	2017	Bangalore	3	34	Male	No	
1	Bachelors	2013	Pune	1	28	Female	No	
2	Bachelors	2014	New Delhi	3	38	Female	No	
3	Masters	2016	Bangalore	3	27	Male	No	
4	Masters	2017	Pune	3	24	Male	Yes	

	ExperienceInCurrentDomain	LeaveOrNot
0	0	0
1	3	1
2	2	0
3	5	1
4	2	1

```
[ ]: missing_data = df.isna()
missing_counts = missing_data.sum()
print(missing_counts)
```

Education	0
JoiningYear	0
City	0
PaymentTier	0
Age	0

```
Gender                0
EverBenched           0
ExperienceInCurrentDomain  0
LeaveOrNot             0
dtype: int64
```

```
[ ]: from sklearn.preprocessing import LabelEncoder
label_encoder = LabelEncoder()
df['Education'] = label_encoder.fit_transform(df['Education'])
df.head()
```

```
[ ]: Education  JoiningYear      City  PaymentTier  Age  Gender  EverBenched  \
0           0         2017  Bangalore           3   34    Male           No
1           0         2013     Pune           1   28   Female           No
2           0         2014  New Delhi           3   38   Female           No
3           1         2016  Bangalore           3   27    Male           No
4           1         2017     Pune           3   24    Male           Yes
```

```
ExperienceInCurrentDomain  LeaveOrNot
0                          0           0
1                          3           1
2                          2           0
3                          5           1
4                          2           1
```

```
[ ]: df['City'] = label_encoder.fit_transform(df['City'])
df.head()
```

```
[ ]: Education  JoiningYear  City  PaymentTier  Age  Gender  EverBenched  \
0           0         2017     0           3   34    Male           No
1           0         2013     2           1   28   Female           No
2           0         2014     1           3   38   Female           No
3           1         2016     0           3   27    Male           No
4           1         2017     2           3   24    Male           Yes
```

```
ExperienceInCurrentDomain  LeaveOrNot
0                          0           0
1                          3           1
2                          2           0
3                          5           1
4                          2           1
```

```
[ ]: df['Gender'] = label_encoder.fit_transform(df['Gender'])
df.head()
```

```
[ ]: Education  JoiningYear  City  PaymentTier  Age  Gender  EverBenched  \
0           0         2017     0           3   34     1           No
```

1	0	2013	2	1	28	0	No
2	0	2014	1	3	38	0	No
3	1	2016	0	3	27	1	No
4	1	2017	2	3	24	1	Yes

	ExperienceInCurrentDomain	LeaveOrNot
0	0	0
1	3	1
2	2	0
3	5	1
4	2	1

```
[ ]: df['EverBenched'] = label_encoder.fit_transform(df['EverBenched'])
df.head()
```

```
[ ]: Education  JoiningYear  City  PaymentTier  Age  Gender  EverBenched  \
0           0         2017     0             3   34      1           0
1           0         2013     2             1   28      0           0
2           0         2014     1             3   38      0           0
3           1         2016     0             3   27      1           0
4           1         2017     2             3   24      1           1
```

	ExperienceInCurrentDomain	LeaveOrNot
0	0	0
1	3	1
2	2	0
3	5	1
4	2	1

```
[ ]: from sklearn.model_selection import train_test_split
```

```
[ ]: X = df.drop('LeaveOrNot',axis=1)
y = df['LeaveOrNot']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4)
```

```
[ ]: from sklearn.metrics import classification_report, confusion_matrix
```

```
[ ]: from sklearn.naive_bayes import GaussianNB
```

```
[ ]: nb_classifier = GaussianNB()
nb_classifier.fit(X_train, y_train)
y_naive_bayes = nb_classifier.predict(X_test)
```

```
[ ]: print("The confusion matrix for Naive Bayes is : ")
print("")
print(confusion_matrix(y_test,y_naive_bayes))
```

The confusion matrix for Naive Bayes is :

```
[[978 238]
 [374 272]]
```

```
[ ]: print("The classification report for Naive Bayes is : ")
      print("")
      print(classification_report(y_test,y_naive_bayes))
```

The classification report for Naive Bayes is :

	precision	recall	f1-score	support
0	0.72	0.80	0.76	1216
1	0.53	0.42	0.47	646
accuracy			0.67	1862
macro avg	0.63	0.61	0.62	1862
weighted avg	0.66	0.67	0.66	1862

```
[ ]: from sklearn.tree import DecisionTreeClassifier
```

```
[ ]: dtree = DecisionTreeClassifier()
      dtree.fit(X_train,y_train)
      y_decision_tree = dtree.predict(X_test)
```

```
[ ]: print("The confusion matrix for Decision Tree is : ")
      print("")
      print(confusion_matrix(y_test,y_decision_tree))
```

The confusion matrix for Decision Tree is :

```
[[1060 156]
 [ 203 443]]
```

```
[ ]: print("The classification report for Decision Tree is : ")
      print("")
      print(classification_report(y_test,y_decision_tree))
```

The classification report for Decision Tree is :

	precision	recall	f1-score	support
0	0.84	0.87	0.86	1216
1	0.74	0.69	0.71	646
accuracy			0.81	1862

macro avg	0.79	0.78	0.78	1862
weighted avg	0.80	0.81	0.81	1862

```
[ ]: from sklearn.ensemble import RandomForestClassifier
```

```
[ ]: rfc = RandomForestClassifier(n_estimators=1000)
rfc.fit(X_train, y_train)
y_random_forest = rfc.predict(X_test)
```

```
[ ]: print("The confusion matrix for Random Forest is : ")
print("")
print(confusion_matrix(y_test, y_random_forest))
```

The confusion matrix for Random Forest is :

```
[[1103  113]
 [ 209  437]]
```

```
[ ]: print("The classification report for Decision Tree is : ")
print("")
print(classification_report(y_test, y_random_forest))
```

The classification report for Decision Tree is :

	precision	recall	f1-score	support
0	0.84	0.91	0.87	1216
1	0.79	0.68	0.73	646
accuracy			0.83	1862
macro avg	0.82	0.79	0.80	1862
weighted avg	0.82	0.83	0.82	1862

```
[ ]: from sklearn.metrics import accuracy_score
```

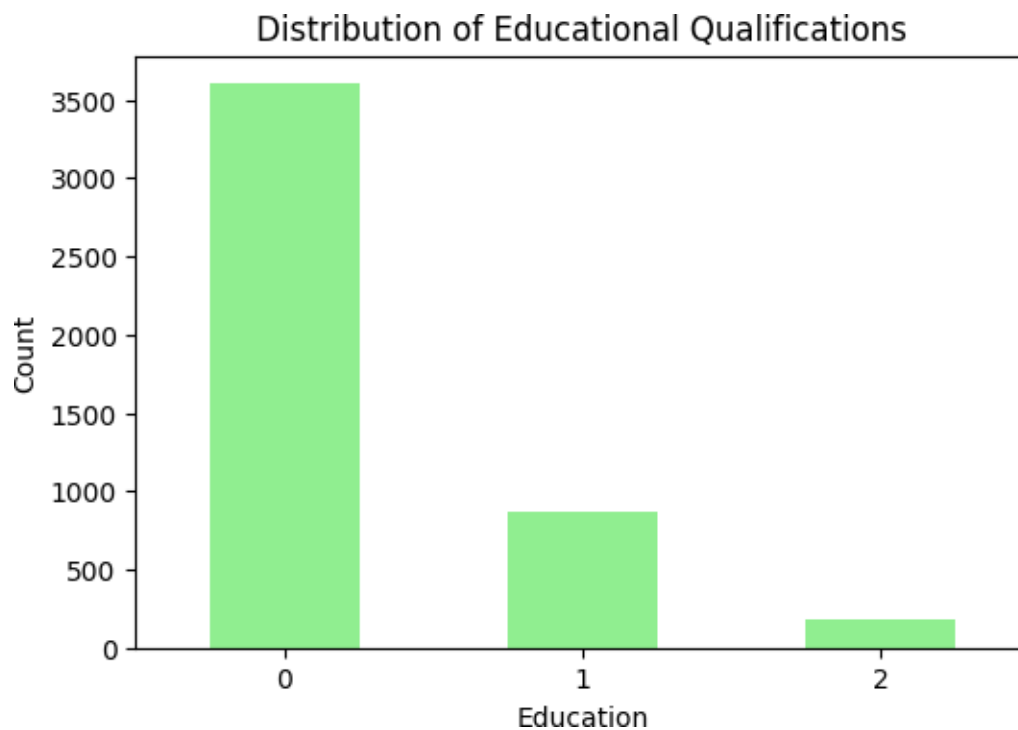
```
[ ]: accuracy_nb = accuracy_score(y_test, y_naive_bayes)
accuracy_dtc = accuracy_score(y_test, y_decision_tree)
accuracy_rfc = accuracy_score(y_test, y_random_forest)

print(f'Accuracy for Naive Bayes: {accuracy_nb:.2f}')
print(f'Accuracy for Decision Tree 1: {accuracy_dtc:.2f}')
print(f'Accuracy for Random Forest Classifier 2: {accuracy_rfc:.2f}')
```

Accuracy for Naive Bayes: 0.67  
Accuracy for Decision Tree 1: 0.81  
Accuracy for Random Forest Classifier 2: 0.83

### Q1. What is the distribution of educational qualifications among employees?

```
[ ]: education_counts = df['Education'].value_counts()
plt.figure(figsize=(6, 4))
education_counts.plot(kind='bar', color='lightgreen')
plt.title('Distribution of Educational Qualifications')
plt.xlabel('Education')
plt.ylabel('Count')
plt.xticks(rotation=0)
plt.show()
```

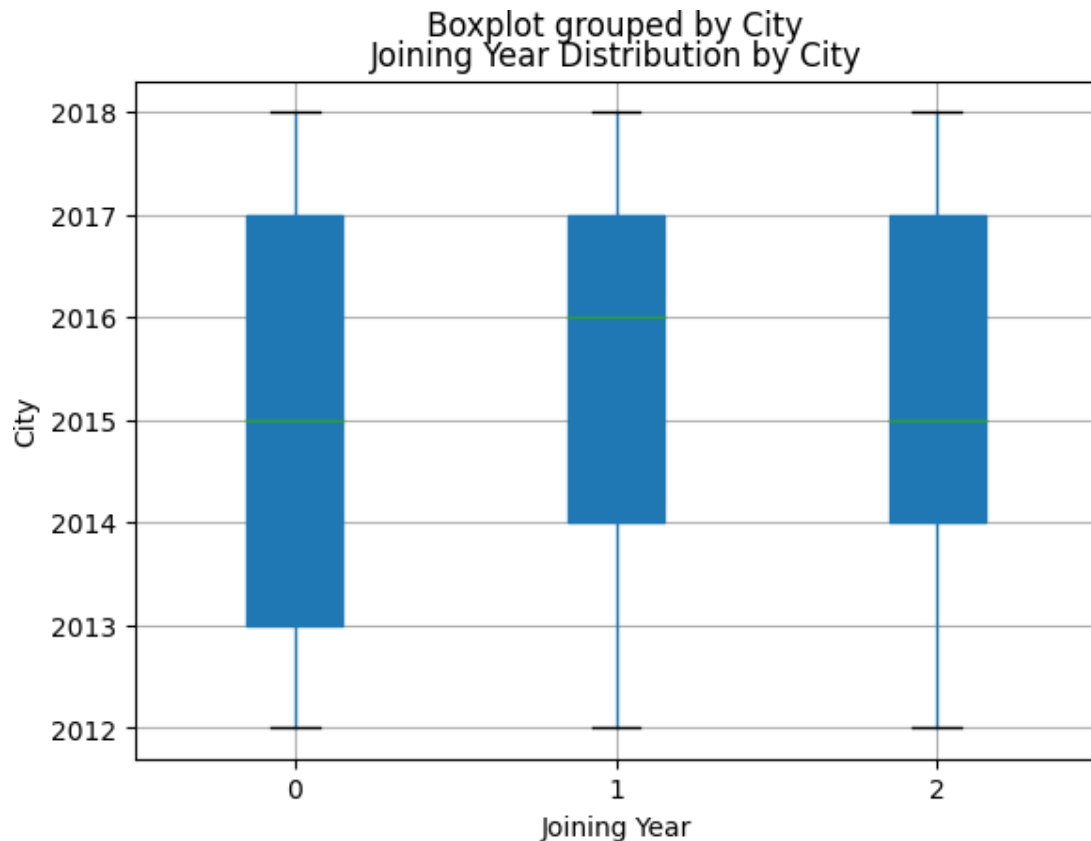


### Q2. How does the length of service (Joining Year) vary across different cities?

```
[ ]: plt.figure(figsize=(10, 6))
df.boxplot(column='JoiningYear', by='City', patch_artist=True)
plt.title('Joining Year Distribution by City')
plt.xlabel('Joining Year')
plt.ylabel('City')

plt.show()
```

<Figure size 1000x600 with 0 Axes>



### Q3. Is there a correlation between Payment Tier and Experience in Current Domain?

```
[ ]: correlation = df['PaymentTier'].corr(df['ExperienceInCurrentDomain'])

print(f"Pearson's Correlation Coefficient: {correlation:.2f}")

if correlation > 0:
    interpretation = "There is a positive correlation."
elif correlation < 0:
    interpretation = "There is a negative correlation."
else:
    interpretation = "There is no linear correlation."


print(interpretation)
```

Pearson's Correlation Coefficient: 0.02  
There is a positive correlation.

### Q4. What is the gender distribution within the workforce?

```
[ ]: gender_counts = df['Gender'].value_counts()
plt.figure(figsize=(6, 4))
gender_counts.plot(kind='bar', color='lightgreen')
plt.title('Gender Distribution in the Workforce')
plt.xlabel('Gender')
plt.ylabel('Count')
plt.xticks(rotation=0)

plt.show()
```

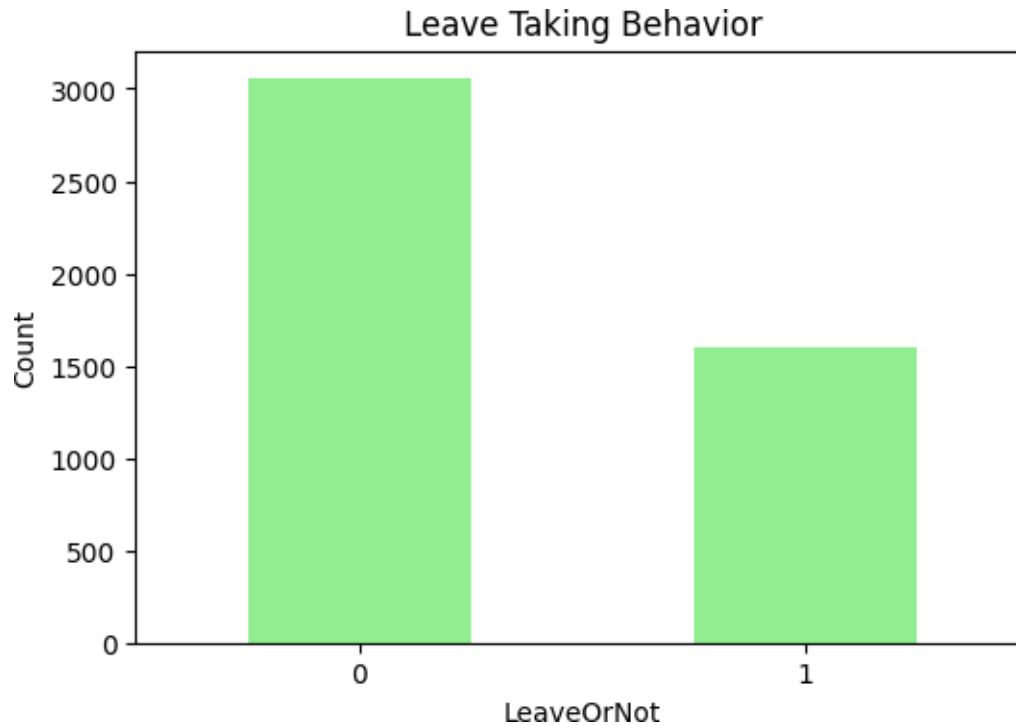


#### Q5. Are there any patterns in leave-taking behavior among employees?

```
[ ]: leave_counts = df['LeaveOrNot'].value_counts()
plt.figure(figsize=(6, 4))
leave_counts.plot(kind='bar', color='lightgreen')
plt.title('Leave Taking Behavior')
plt.xlabel('LeaveOrNot')
plt.ylabel('Count')
plt.xticks(rotation=0)

plt.show()
```





```
[ ]: correlation_leave = df.corr()['LeaveOrNot'].drop('LeaveOrNot')
      print(correlation_leave)
```

```
Education          0.080497
JoiningYear        0.181705
City               0.201058
PaymentTier       -0.197638
Age               -0.051126
Gender             -0.220701
EverBenched        0.078438
ExperienceInCurrentDomain -0.030504
Name: LeaveOrNot, dtype: float64
```

```
[ ]: plt.scatter(df['ExperienceInCurrentDomain'], df['LeaveOrNot'], alpha=0.5)
      plt.title('Experience in Current Domain vs. Leave Taking')
      plt.xlabel('Experience in Current Domain')
      plt.ylabel('LeaveOrNot')
      plt.show()

      df.boxplot(column='PaymentTier', by='LeaveOrNot', patch_artist=True)
      plt.title('PaymentTier by Leave Taking Behavior')
      plt.xlabel('PaymentTier')
      plt.ylabel('LeaveOrNot')
```

```
plt.show()
```

